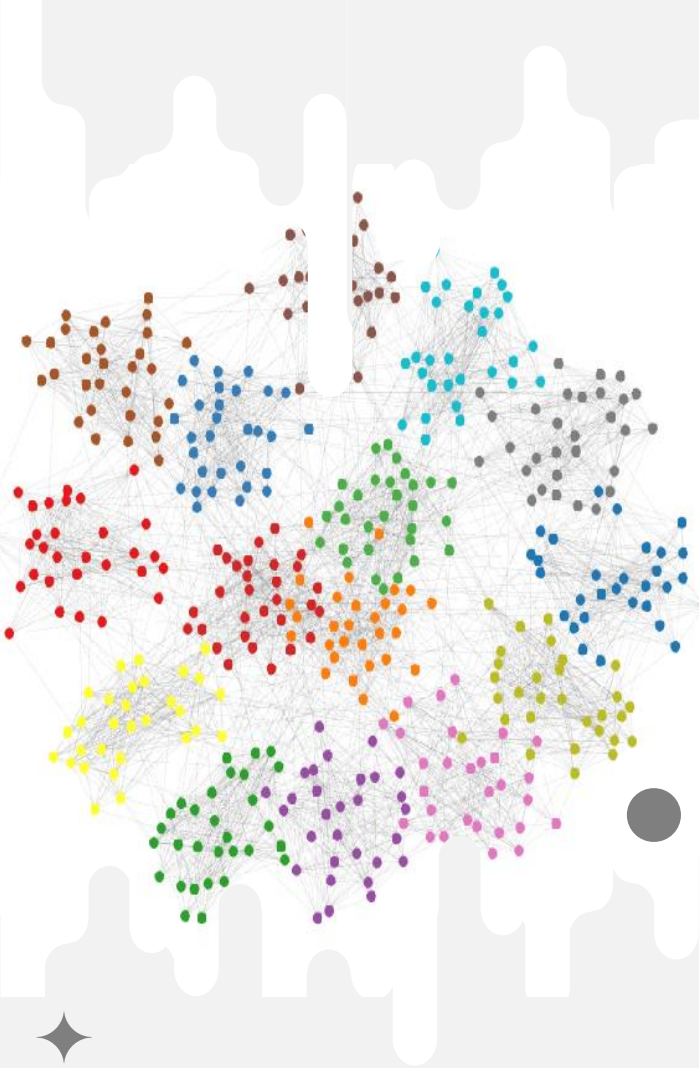




Model Compression Homework

繳交期限：2025 / 05 / 27 (二) 23:59

TA email: lobsterlab.cs.nthu@gmail.com



大綱

We are going to cover ...

01

作業介紹

02

提供的 **code**

03

繳交內容

04

配分標準



1

作業介紹

DEEP COMPRESSION

Python library you need

DEEP COMPRESSION


- 架構：



Prune


DEEP COMPRESSION

- 架構：




DEEP COMPRESSION

- 架構：



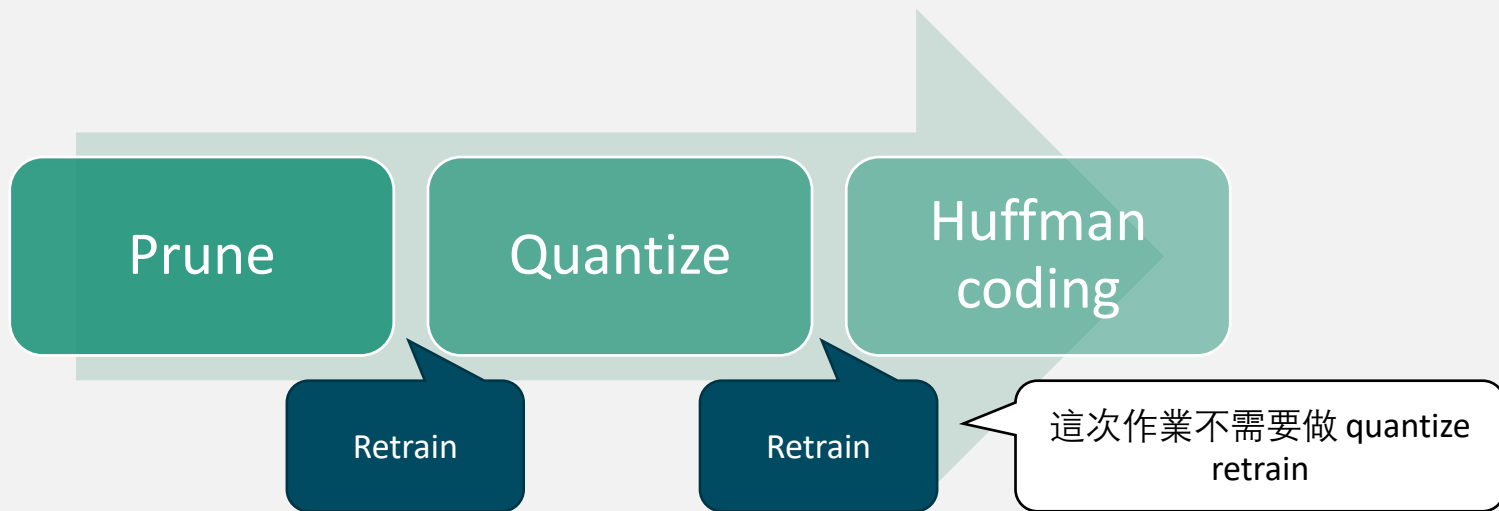
DEEP COMPRESSION

- 架構：



DEEP COMPRESSION

- 架構：



Python library

- pytorch
 - 深度學習套件
 - 記得載 GPU 版



Python library

- torchvision
 - 下載常用測試 data 套件
 - 本次資料使用 cifar10
 - 也可用來下載mnist、cifar100



Python library

- scikit-learn (sklearn)
 - 機器學習套件
 - 本次要使用 cluster 的演算法 : Kmeans
 - (在 quantize 那一步)



2

提供的 code



模型



Prune 部分

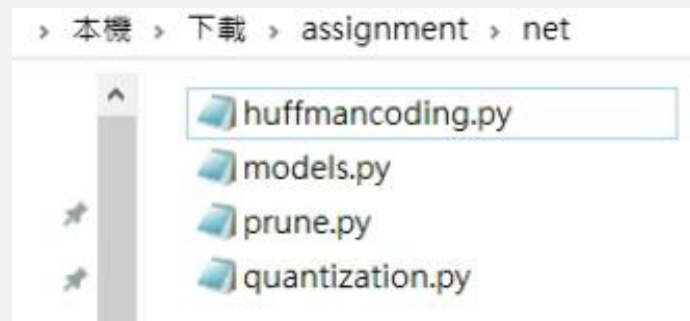
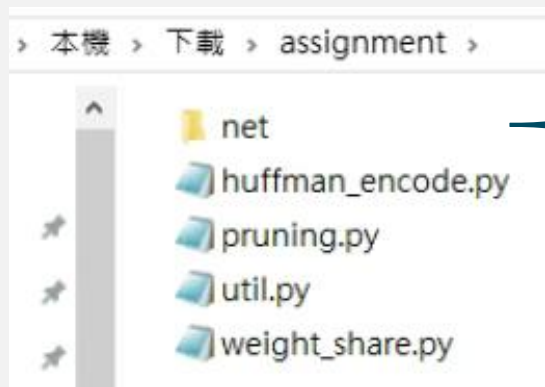


Quantize 部分



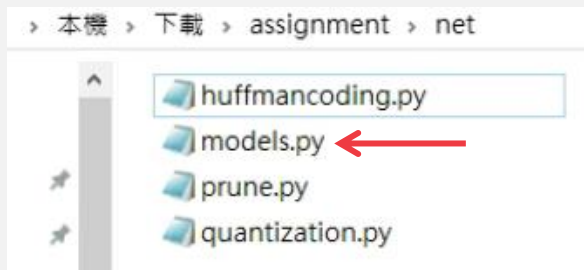
Huffman 部分

Assignment



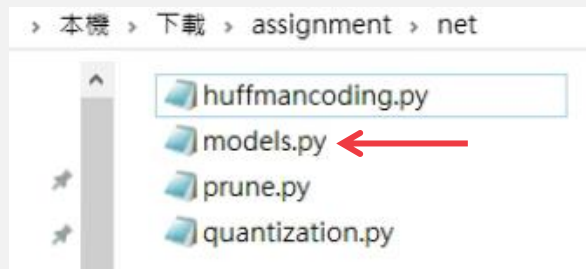
模型

- AlexNet



模型

- AlexNet



```
class AlexNet(PruningModule):
    def __init__(self, n_classes=10):
        super(AlexNet, self).__init__()

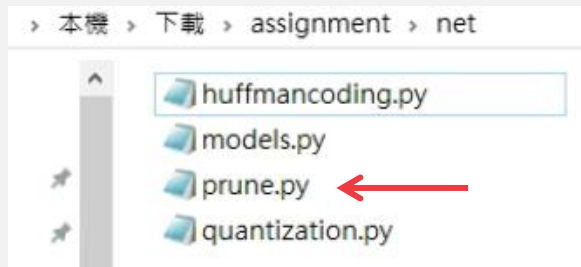
        self.conv1 = nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=5)
        self.conv2 = nn.Conv2d(64, 192, kernel_size=5, padding=2)
        self.conv3 = nn.Conv2d(192, 384, kernel_size=3, padding=1)
        self.conv4 = nn.Conv2d(384, 256, kernel_size=3, padding=1)
        self.conv5 = nn.Conv2d(256, 256, kernel_size=3, padding=1)

        self.fc1 = nn.Linear(256, 4096)
        self.fc2 = nn.Linear(4096, 4096)
        self.fc3 = nn.Linear(4096, n_classes)

    def forward(self, x):
        x = F._max_pool2d(F.relu(self.conv1(x), inplace=True), kernel_size=2)
        x = F._max_pool2d(F.relu(self.conv2(x), inplace=True), kernel_size=2)
        x = F.relu(self.conv3(x), inplace=True)
        x = F.relu(self.conv4(x), inplace=True)
        x = F._max_pool2d(F.relu(self.conv5(x), inplace=True), kernel_size=2)
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x), inplace=True)
        x = F.relu(self.fc2(x), inplace=True)
        x = self.fc3(x)
        x = F.log_softmax(x, dim=1)
        return x
```

Prune

- prune.py
- 只提供 prune by std 的 fully connected 部分 (壓縮率不可預期)
- 可選擇實作 prune by percentile 與否 (壓縮率可預期)



prune.py

```
def prune_by_std(self, s=0.25):
    for name, module in self.named_modules():

        #####
        # TODO:
        #    Only fully connected layers were considered, but convolution layers also needed
        #####

        if name in ['fc1', 'fc2', 'fc3']:
            threshold = np.std(module.weight.data.cpu().numpy()) * s
            print(f'Pruning with threshold : {threshold:.4f} for layer {name}')
            self._prune(module, threshold)
```

prune.py

```
def _prune(self, module, threshold):  
  
    #####  
    # TODO:  
    # 1. Use "module.weight.data" to get the weights of a certain layer of the model  
    # 2. Set weights whose absolute value is less than threshold to 0, and keep the rest unchanged  
    # 3. Save the results of the step 2 back to "module.weight.data"  
    # -----  
    # In addition, there is no need to return in this function ("module" can be considered as call by  
    # reference)  
    #####  
  
    pass
```

prune.py

```
DEFAULT_PRUNE_RATE = {
    'conv1': 84,
    'conv2': 38,
    'conv3': 35,
    'conv4': 37,
    'conv5': 37,
    'fc1': 9,
    'fc2': 9,
    'fc3': 25
}

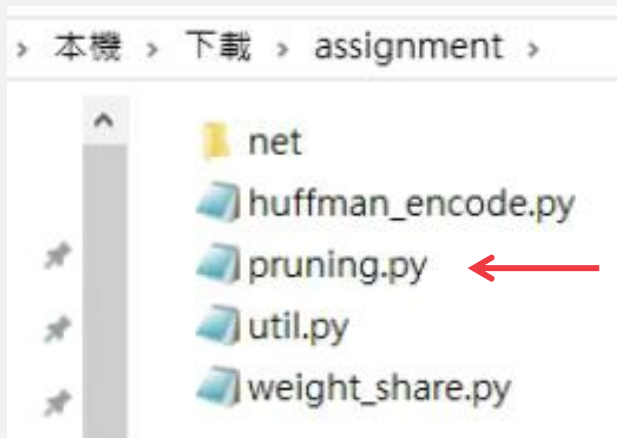
def prune_by_percentile(self, q=DEFAULT_PRUNE_RATE):
    #####
    # TODO
    # For each layer of weights W (including fc and conv layers) in the model, obtain the (100 - q)th percentile
    # of absolute W as the threshold, and then set the absolute weights less than threshold to 0, and the rest
    # remain unchanged.
    #####

    # Calculate percentile value
    pass

    # Prune the weights and mask
    pass
```

Prune

- pruning.py
- Initial train + prune + prune retrain 的部分



pruning.py

```
def train(epochs):
    model.train()
    for epoch in range(epochs):
        pbar = tqdm(enumerate(train_loader), total=len(train_loader))
        for batch_idx, (data, target) in pbar:
            data, target = data.to(device), target.to(device)
            optimizer.zero_grad()
            output = model(data)
            loss = F.nll_loss(output, target)
            loss.backward()

            for name, p in model.named_parameters():

                #####
                # TODO:
                #   zero-out all the gradients corresponding to the pruned weights
                #####

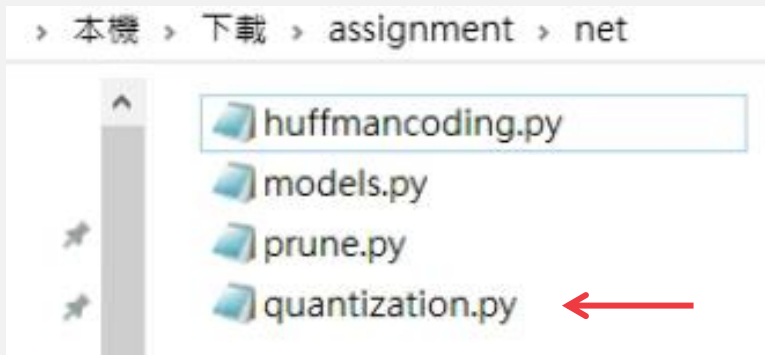
            pass
```

Run : pruning.py

- 生出 :  model_after_retraining.ptmodel

Quantize

- quantization.py
- 只提供 fully connected 部分



quantization.py

```
def apply_weight_sharing(model, bits=5):
    """
    Applies weight sharing to the given model
    """
    for name, module in model.named_children():
        dev = module.weight.device
        weight = module.weight.data.cpu().numpy()
        shape = weight.shape
        quan_range = 2 ** bits
        if len(shape) == 2: # Fully connected layers
            print(f'{name:20} | {str(module.weight.size()):35} | => Quantize to {quan_range} indices')
            mat = csr_matrix(weight) if shape[0] < shape[1] else csc_matrix(weight)

            # Weight sharing by kmeans
            space = np.linspace(min(mat.data), max(mat.data), num=quan_range)
            kmeans = KMeans(
                n_clusters=len(space),
                init=space.reshape(-1, 1),
                n_init=1,
                precompute_distances=True,
                algorithm="full"
            )
            kmeans.fit(mat.data.reshape(-1, 1))
            new_weight = kmeans.cluster_centers_[kmeans.labels_].reshape(-1)
            mat.data = new_weight

        # Insert to model
        module.weight.data = torch.from_numpy(mat.toarray()).to(dev)
```


quantization.py

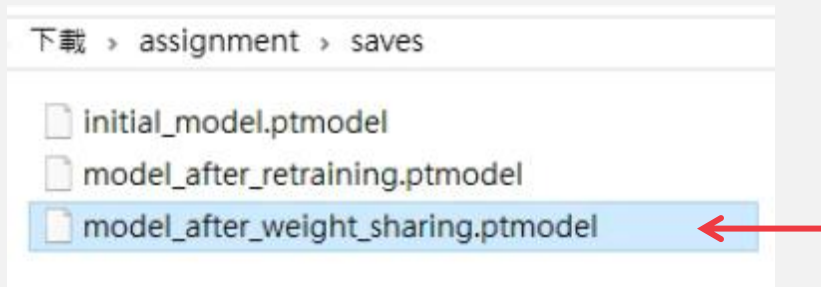
```
elif len(shape) == 4: # Convolution layers

#####
# TODO:
#   Suppose the weights of a certain convolution layer are called "W"
#   1. Get the unpruned (non-zero) weights, "non-zero-W", from "W"
#   2. Use KMeans algorithm to cluster "non-zero-W" to (2 ** bits) categories
#   3. For weights belonging to a certain category, replace their weights with the centroid
#       value of that category
#   4. Save the replaced weights in "module.weight.data", and need to make sure their indices
#       are consistent with the original
#   Finally, the weights of a certain convolution layer will only be composed of (2 ** bits) float numbers
#   and zero
#   -----
#   In addition, there is no need to return in this function ("model" can be considered as call by
#   reference)
#####

print(f'{name:20} | {str(module.weight.size()):35} | ** NEED TO BE IMPLEMENTED **')
pass
```

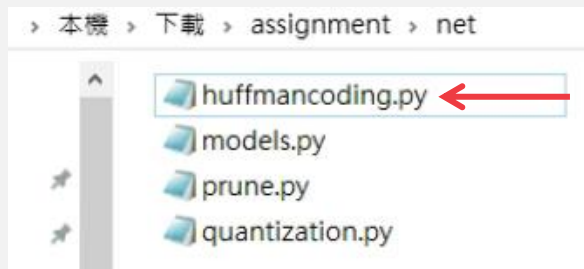
Run : weight_share.py

- 生出 :



Huffman coding

- huffmancoding.py
- 分別有 encode、decode 部分
- 且都只提供 fully connected 部分



Huffman encode

```
# Encode / Decode models
def huffman_encode_model(model, directory='encodings/'):
    os.makedirs(directory, exist_ok=True)
    original_sum = 0
    compressed_sum = 0
    log_text = f'{"Layer":<15} | {"original bytes":>20} {"compressed bytes":>20} {"improvement":>11} {"percent":>7}\n'
    log_text += '-' * 70
    util.log(log_text)
    print(log_text)
    for name, param in model.named_parameters():
        if 'weight' in name: # Weights
            if 'conv' in name:
                original, compressed = huffman_encode_conv(param, name, directory)
            elif 'fc' in name:
                original, compressed = huffman_encode_fc(param, name, directory)
            else:
                raise NameError
        else: # Bias
            original, compressed = dump_bias(param, name, directory)
        original_sum += original
        compressed_sum += compressed
    log_text = '-' * 70 + '\n'
    log_text += (
        f'{"total":15} | {"original_sum":>20} {"compressed_sum":>20} {"original_sum / compressed_sum":>10.2f}x "
        f"{100 * compressed_sum / original_sum:>6.2f}%"
    )
    util.log(log_text)
    print(log_text)
```

```

def huffman_encode_fc(param, name, directory):
    weight = param.data.cpu().numpy()
    shape = weight.shape

    form = 'csr' if shape[0] < shape[1] else 'csc'
    mat = csr_matrix(weight) if shape[0] < shape[1] else csc_matrix(weight)

    # Encode
    t0, d0 = huffman_encode(mat.data, name + f'_{form}_data', directory)
    t1, d1 = huffman_encode(mat.indices, name + f'_{form}_indices', directory)
    t2, d2 = huffman_encode(calc_index_diff(mat.indptr), name + f'_{form}_indptr', directory)

    # Print statistics
    original = param.data.cpu().numpy().nbytes
    compressed = t0 + t1 + t2 + d0 + d1 + d2
    log_text = (
        f"{name:<15} | {original:20} {compressed:20} {original / compressed:>10.2f}x "
        f"{100 * compressed / original:>6.2f}%"
    )
    util.log(log_text)
    print(log_text)

    return original, compressed

```

```

def huffman_encode_conv(param, name, directory):
    #####
    # TODO:
    # You can refer to the code of the function "huffman_encode_fc" below, but note that "csr_matrix" can only be
    # used on 2-dimensional data
    # -----
    # HINT:
    # Suppose the shape of the weights of a certain convolution layer is (Kn, Ch, W, H)
    # ---
    # 1. Call function "csr_matrix" for all (Kn * Ch) two-dimensional matrices (W, H), and get "data",
    # "length of data", "indices", and "indptr" of all (Kn * Ch) csr_matrix.
    # 2. Concatenate these 4 parts of all (Kn * Ch) csr_matrices individually into 4 one-dimensional
    # lists, so there will be 4 lists.
    # 3. Do huffman coding on these 4 lists individually.
    #####

    # Note that we do not huffman encode "conv" yet. The following four lines of code need to be modified
    conv = param.data.cpu().numpy()
    conv.dump(f'{directory}/{name}')

    # Print statistics
    original = conv.nbytes
    compressed = original
    log_text = (
        f'{name:<15} | '
        f'{original:20} {compressed:20} {original / compressed:>10.2f}x '
        f'{100 * compressed / original:>6.2f}% (NEED TO BE IMPLEMENTED)'
    )
    util.log(log_text)
    print(log_text)

    return original, compressed

```

Huffman decode

```
def huffman_decode_model(model, directory='encodings/'):
    for name, param in model.named_parameters():
        if 'weight' in name:
            if 'conv' in name:
                huffman_decode_conv(param, name, directory)
            elif 'fc' in name:
                huffman_decode_fc(param, name, directory)
            else:
                raise NameError
        else:
            load_bias(param, name, directory)
```

```

def huffman_decode_conv(param, name, directory):
    #####
    # TODO:
    # Decode according to the code of "conv" section you write in the function "huffman encode model"
    # above, and refer to encode and decode code of "fc"
    #####

    # Note that we do not huffman decode "conv" yet. The following three lines of code need to be modified
    conv = np.load(directory + '/' + name, allow_pickle=True)
    param.data = torch.from_numpy(conv).to(param.device)

def huffman_decode_fc(param, name, directory):
    weight = param.data.cpu().numpy()
    shape = weight.shape

    form = 'csr' if shape[0] < shape[1] else 'csc'
    matrix = csr_matrix if shape[0] < shape[1] else csc_matrix

    # Decode data
    data = huffman_decode(directory, name + f'_{form}_data', dtype='float32')
    indices = huffman_decode(directory, name + f'_{form}_indices', dtype='int32')
    indptr = reconstruct_indptr(huffman_decode(directory, name + f'_{form}_indptr', dtype='int32'))

    # Construct matrix
    mat = matrix((data, indices, indptr), shape)

    # Insert to model
    param.data = torch.from_numpy(mat.toarray()).to(param.device)

```


Run : huffman_coding.py

- 觀察結果：
 - 壓縮率
 - accuracy

```
--- Start encoding ---
Layer | original bytes compressed bytes improvement percent
-----|-----
conv1.weight | 92928 2457 37.82x 2.64%
conv1.bias | 256 256 1.00x 100.00%
conv2.weight | 1228800 34336 35.79x 2.79%
conv2.bias | 768 768 1.00x 100.00%
conv3.weight | 2654208 79993 33.18x 3.01%
conv3.bias | 1536 1536 1.00x 100.00%
conv4.weight | 3538944 91339 38.75x 2.58%
conv4.bias | 1024 1024 1.00x 100.00%
conv5.weight | 2359296 53404 44.18x 2.26%
conv5.bias | 1024 1024 1.00x 100.00%
fc1.weight | 4194304 44834 93.55x 1.07%
fc1.bias | 16384 16384 1.00x 100.00%
fc2.weight | 67108864 138973 482.89x 0.21%
fc2.bias | 16384 16384 1.00x 100.00%
fc3.weight | 163840 2732 59.97x 1.67%
fc3.bias | 40 40 1.00x 100.00%
-----|-----
total | 81378600 485484 167.62x 0.60%
--- Start decoding ---
--- Accuracy after decoding ---
Test set: Average loss: 2.6773, Accuracy: 5642/10000 (56.42%)
```



3

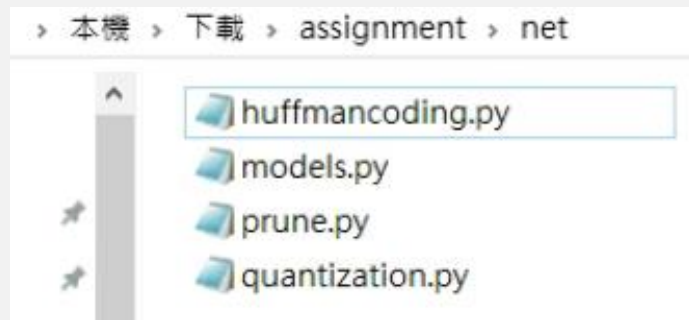
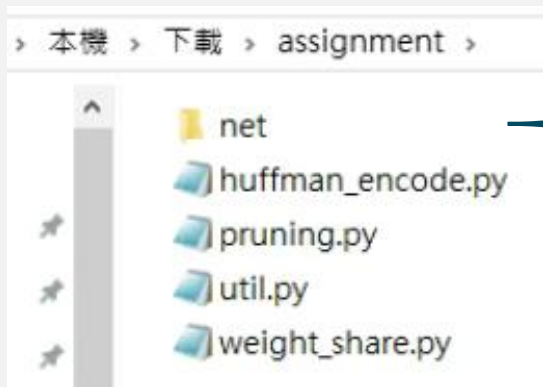
繳交內容

TODO

- 有 TODO 的 py 檔：
 - prune.py
 - pruning.py
 - quantization.py
 - huffmancoding.py
- 請根據 TODO 的提示內容完成程式

繳交規格

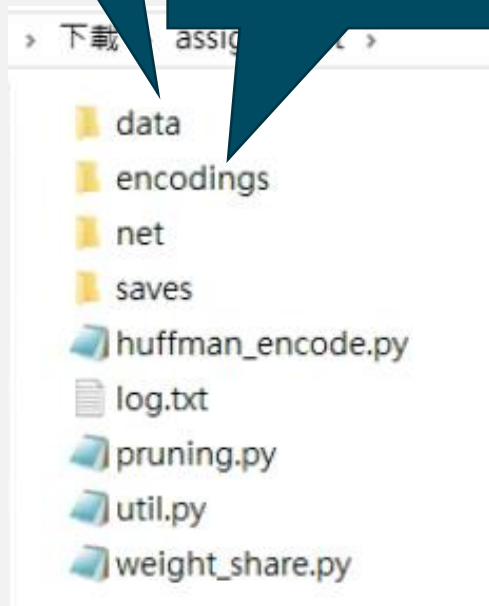
- 原本拿到的



繳交規格

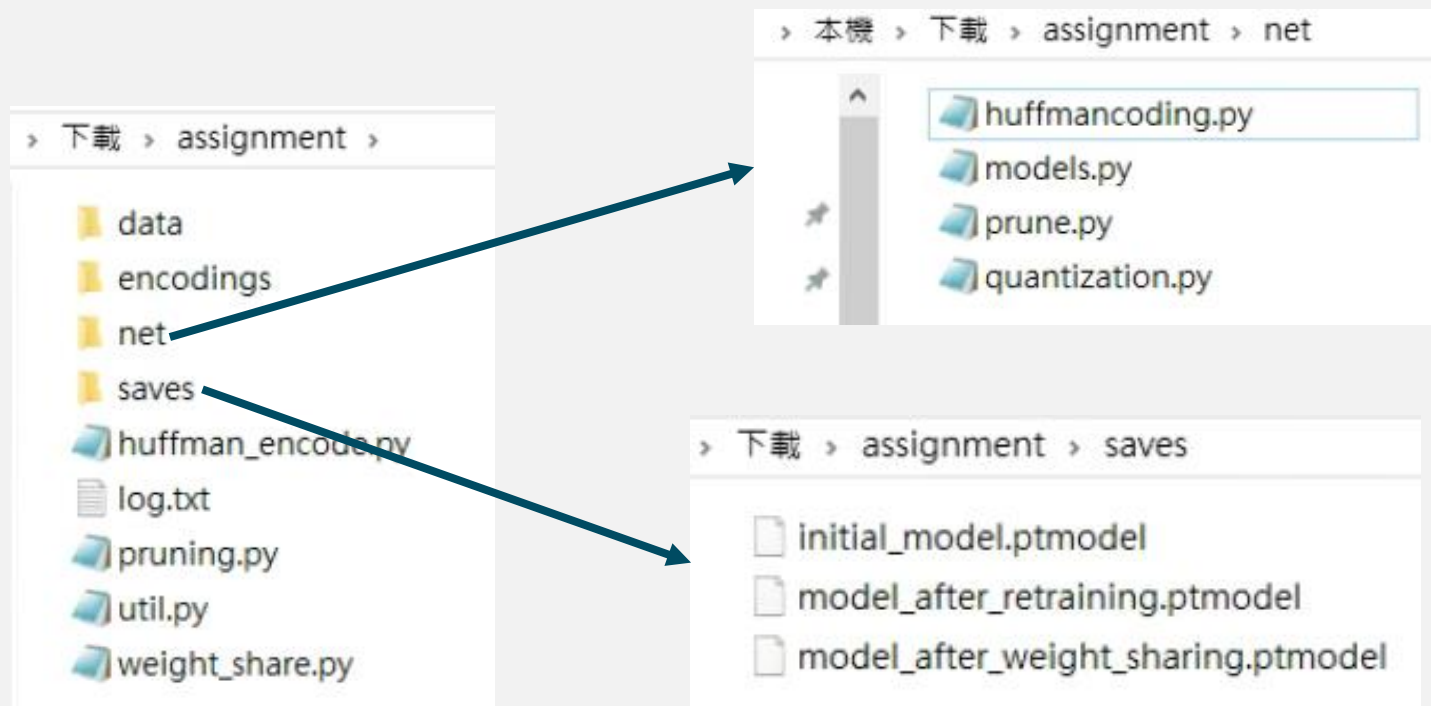
torchvision 下載的 data

huffmancoding 所生成


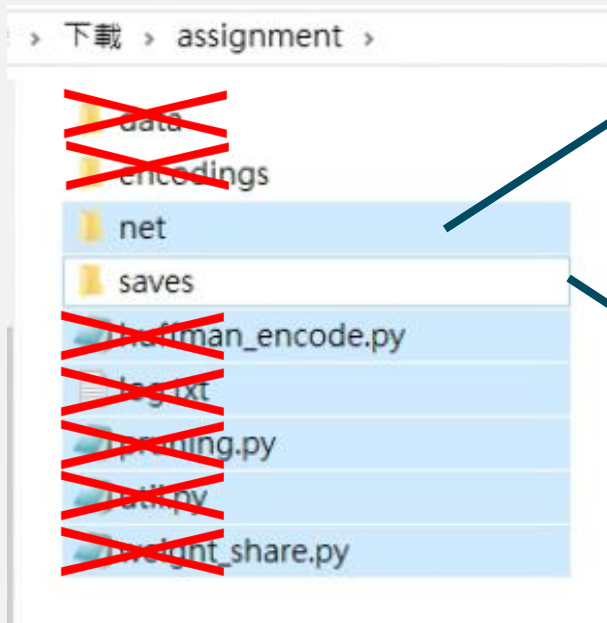


- 跑完程式後，應該長這樣

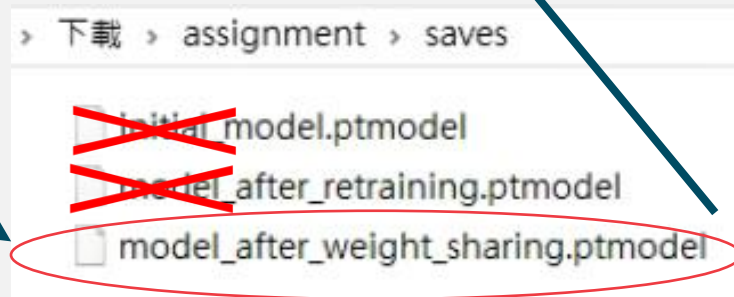
繳交規格



繳交規格



繳交格式: 學號.rar、學號.ptmodel



繳交規格



108062345.rar



108062345.ptmodel

最終繳交這兩項

4

配分標準

配分

- (70%) Prune model : 完成 `prune.py` 、 `pruning.py` 中 # TODO
- (10%) Quantize model : 完成 `quantization.py` 中 # TODO
- (10%) Huffman coding : 完成 `huffmancoding.py` 中 # TODO
- (10%) 壓縮率排名 : $\text{accuracy} > 58\%$, 最終模型 `compression rate` 越高越好 ! 排名前1/3得10分、中1/3得5分、後1/3得0分 , 會用同學繳回的模型計算 , 排名改完會公布於 `eeclass`



Thank you

鼓勵討論，嚴禁抄襲